

Assignment -2

Data Visualization and Pre-processing

Assignment Date	30 September 2022
Student Name	R.Roja
Student Roll Number	CS19035
Maximum Marks	2 Marks

In []:

```
# Importing required libraries
```

```
import numpy as np
import pandas as pd
```

In []:

```
# Reading the dataset
```

```
df = pd.read_csv('/content/Churn_Modelling.csv')
```

In []:

```
# Visualizing 1st 50 data
```

```
df.head()
```

Out[]:

	0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1
RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCar	
	1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	

In []:

Checking for null values

df.isnull().sum() Out[]:

RowNumber	0
CustomerId	0
Surname	0
CreditScore	0
Geography	0
Gender	0
Age	0
Tenure	0
Balance	0
NumOfProducts	0
HasCrCard	0
IsActiveMember	0
EstimatedSalary	0
Exited dtype:	0
int64	

In []:

df.dtypes

Out[]:

RowNumber	int64
CustomerId	int64
Surname	object
CreditScore	int64
Geography	object
Gender	object
Age	int64
Tenure	int64
Balance	float64
NumOfProducts	int64
HasCrCard	int64
IsActiveMember	int64
EstimatedSalary	float64
Exited	

int64 dtype:

object

```
import matplotlib.pyplot as plt
import seaborn as sns
```

In [47]:



Univariate Analysis

```
sns.histplot(data["CreditScore"], color='darkorange')
```

In [48]:

Out[48]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f831677f6d0>



```
sns.displot(data['Age'], kde=True)
```

In [49]:

Out[49]:

<seaborn.axisgrid.FacetGrid at 0x7f831661b210>


Bi - Variate Analysis

```
sns.barplot(data=data.head(50), x="Geography", y="CreditScore")
```

In [50]:

Out[50]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f8313ce63d0>




```
sns.catplot(x='Gender', y='CreditScore', hue='HasCrCard', data=data)
```

In [51]:

Out[51]:

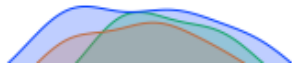
<seaborn.axisgrid.FacetGrid at 0x7f8317198a90>

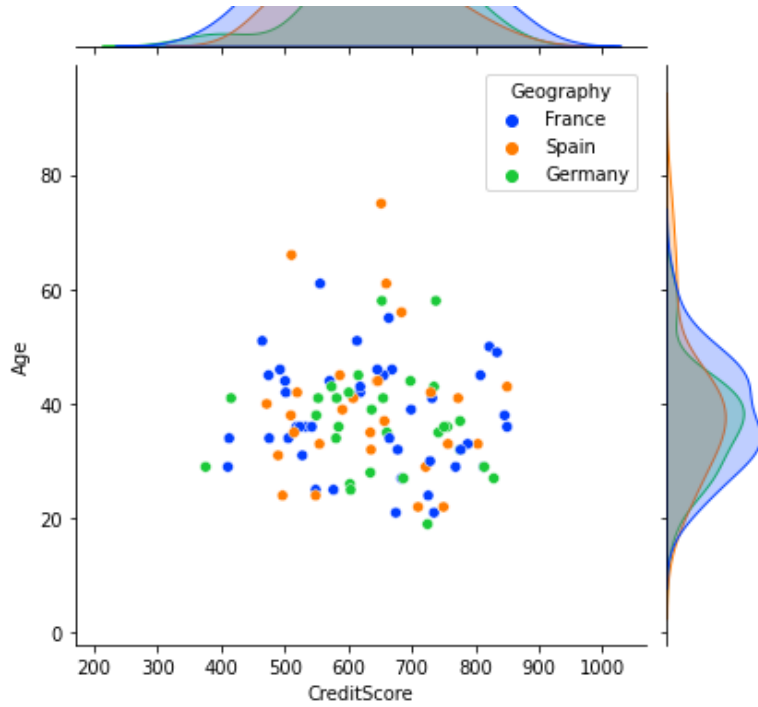


Multi - Variate Analysis

```
sns.jointplot(  
    x='CreditScore',  
    y='Age',  
    data=data.head(100),  
    palette='bright',  
    hue='Geography');
```

In [52]:



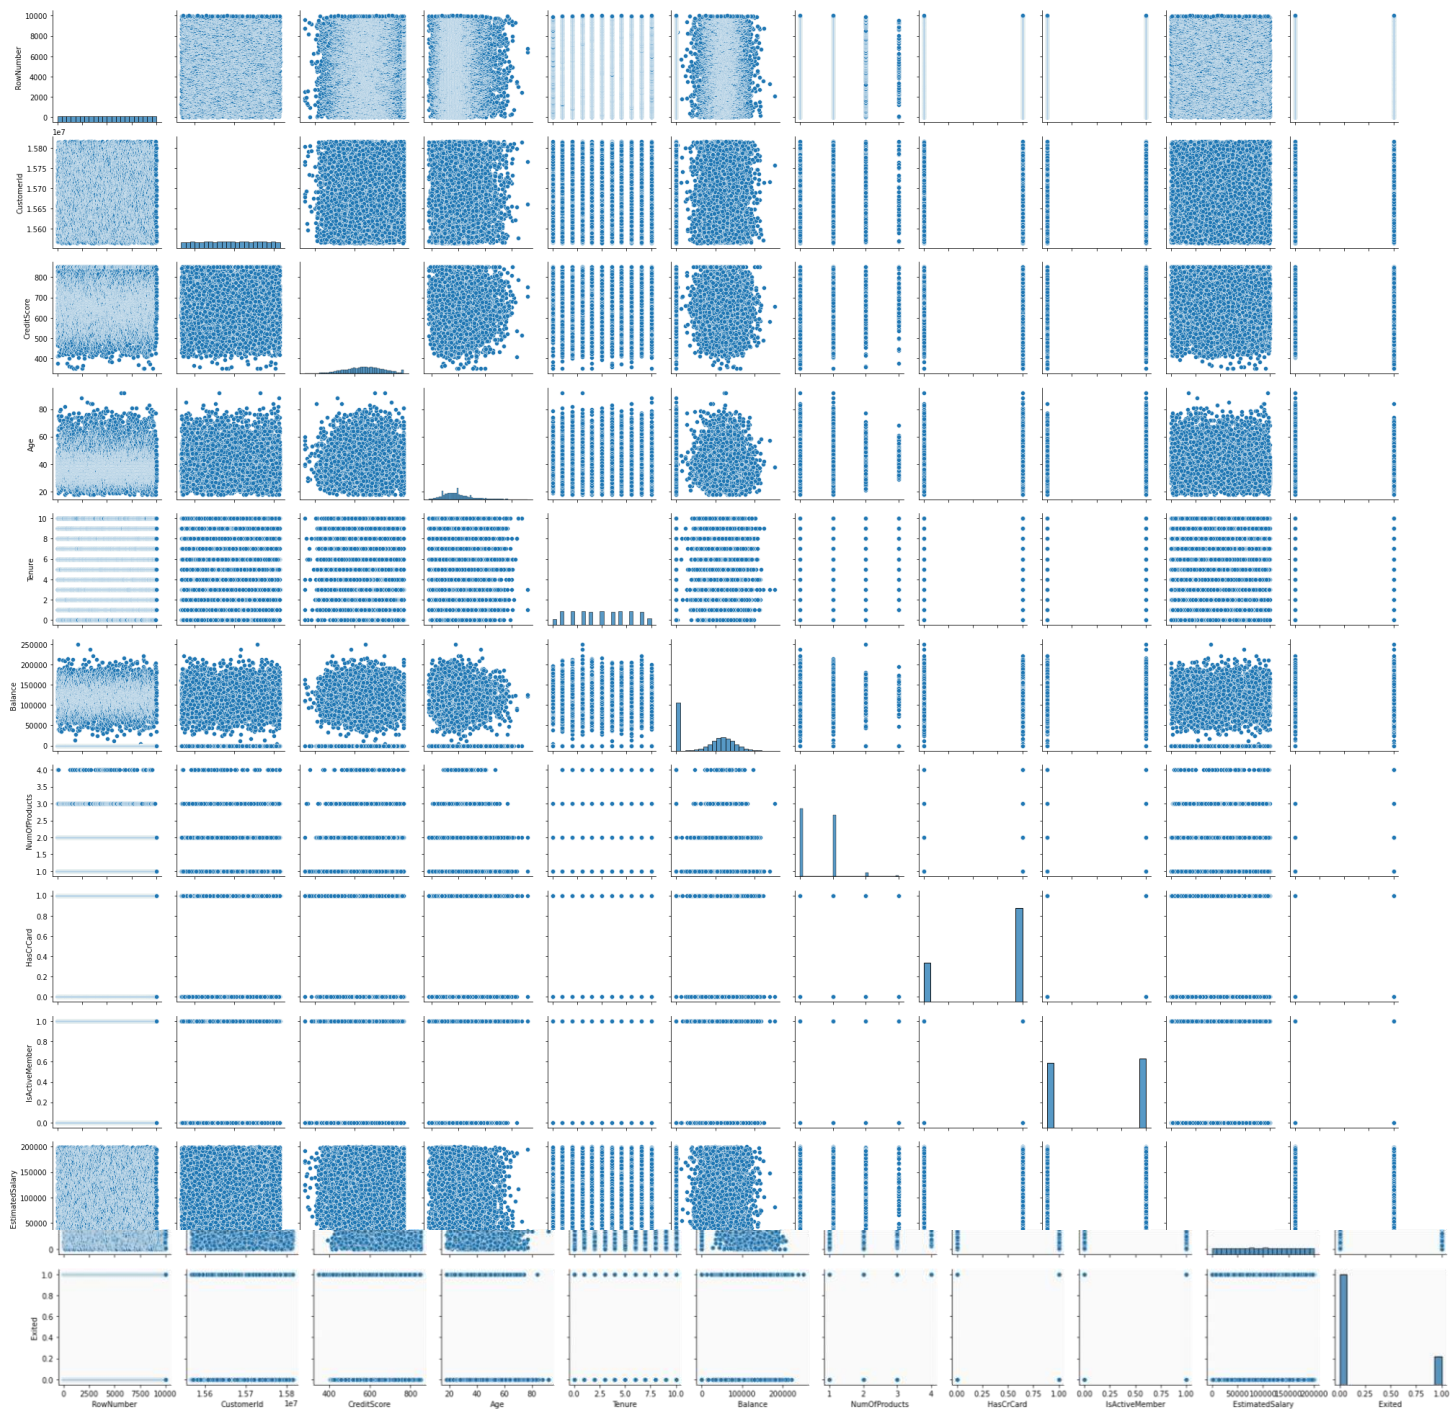


```
sns.pairplot(data)
```

In [53]:

Out[53]:

<seaborn.axisgrid.PairGrid at 0x7f8313a71390>



Perform descriptive statistics on the dataset

```
data.describe()
```

In [54]:

Out[54]:

	RowNum	CustomerID	CreditScore	Age	Tenure	Balance	NumOfProd	HasCrCa	InterestRate	Products
count	10000.000	1.000000e+07	10000.000	10000.000	10000.000	10000.000	10000.0000	10000.00000		10000.00000
mean	5000.5000	1.569094e+07	650.528800	38.921800	5.012800	76485.8892		1.530200		0.70550
std	2886.8956	7.193619e+04	96.653299		10.487806	2.892174	62397.4052	0.581654		0.45584
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000		0.000000	1.000000		0.00000
25%	2500.7500	1.562853e+07	584.000000	32.000000	3.000000		0.000000	1.000000		0.00000
50%	5000.5000	1.569074e+07	652.000000	37.000000	5.000000	97198.5400		1.000000		1.00000
75%	7500.2500	1.575323e+07	718.000000	44.000000	7.000000	127644.240000		2.000000		1.00000
max	10000.000	1.581569e+07	850.000000	92.000000	10.000000	250898.090	4.000000			1.00000

Handle the Missing values

```
data.isnull().sum()
```

In [55]:

Out[55]:

RowNumber	0
CustomerId	0
Surname	0
CreditScore	0
Geography	0
Gender	0
Age	0
Tenure	0
Balance	0
NumOfProducts	0
HasCrCard	0
IsActiveMember	0
EstimatedSalary	0
Exited	0

dtype: int64

Find the outliers and replace the outliers

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result

```
import seaborn as sns

sns.boxplot(data['CreditScore'])
```

in an error or misinterpretation.

In [56]:

FutureWarning

Out[56]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f8310b82990>

```

import numpy as np

Q1 = np.percentile(data['CreditScore'], 25,
                    interpolation = 'midpoint')

Q3 = np.percentile(data['CreditScore'], 75,
                    interpolation = 'midpoint')
IQR = Q3 - Q1

#Upper bound
upper = np.where(data['CreditScore'] >= (Q3+1.5*IQR))
#Lower bound
lower = np.where(data['CreditScore'] <= (Q1-1.5*IQR))

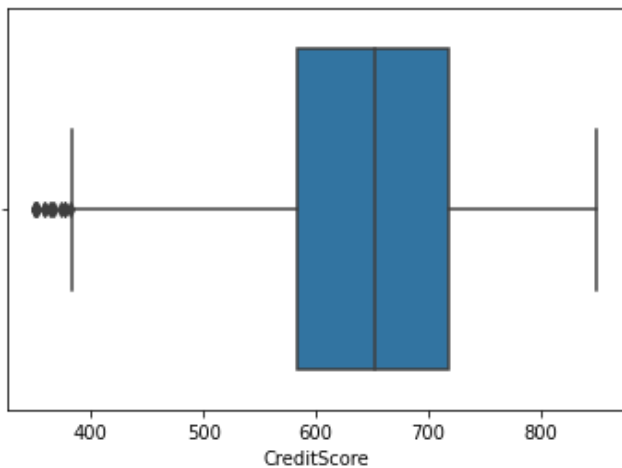
print("Q3: ",Q3)
print("Q1: ",Q1)
print("IQR: ",IQR)

mean = data["CreditScore"].mean()

data["CreditScore"] = np.where(data["CreditScore"] > 850, mean, data['CreditScore'])
data["CreditScore"] = np.where(data["CreditScore"] < 400, mean, data['CreditScore'])

sns.boxplot(data['CreditScore'])

```



In [57]:

Q3: 718.0


Q1: 584.0

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation. FutureWarning

IQR: 134.0 Out[57]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f83177a7310>





Check for Categorical columns and perform encoding

```
from sklearn.preprocessing import LabelEncoder

le = LabelEncoder()
data['Geography'] = le.fit_transform(data['Geography'])
data['Gender'] = le.fit_transform(data['Gender'])

data.head()
```

In [58]:

Out[58]:

RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance
0	15634602	Hargrave	619.0	0	0	42	2	0.00
1	15647311	Hill	608.0	2	0	41	1	83807.86
2	15619304	Onio	502.0	0	0	42	8	159660.80
4	15737888	Mitchell	850.0	2	0	43	2	125510.82
3	15701354	Boni	699.0	0	0	39	1	0.00

Split the data into dependent and independent variables

```
y = data['CreditScore'] #dependent
x = data.drop(columns = ['CreditScore'],axis = 1) #independent

x . head()
```

In [59]:

Out[59]:

	RowNumber	Customer Id	Surname	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
0	1	15634602	Hargrave	0	0	42	2	0.00	1	1	
1	2	15647311	Hill	2	0	41	1	83807.86	1	0	
2	3	15619304	Onio	0	0	42	8	159660.80	3	1	
3	4	15701354	Boni	0	0	39	1	0.00	2	0	
4	5	15737888	Mitchell	2	0	43	2	125510.82	1	1	



Scale the independent variables

```
names = ['RowNumber', 'CustomerId', 'Geography', 'Gender', 'Age', 'Tenure', 'Balance', 'NumOfPro
```

In [60]:

```
ducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary', 'Exited']
```

In [61]:

```
from sklearn.preprocessing import scale
```

```
x = scale(x[names]) x Out[61]:
```

```
array([[ -1.73187761, -0.78321342, -0.90188624, ...,  0.97024255,
         0.02188649,  1.97716468],
       [ -1.7315312 , -0.60653412,  1.51506738, ...,  0.97024255,
         0.21653375, -0.50577476],
       [ -1.73118479, -0.99588476, -0.90188624, ..., -1.03067011,
         0.2406869 ,  1.97716468],
       ...,
       [  1.73118479, -1.47928179, -0.90188624, ...,  0.97024255,
        -1.00864308,  1.97716468],
       [  1.7315312 , -0.11935577,  0.30659057, ..., -1.03067011,
        -0.12523071,  1.97716468],
```

```
[ 1.73187761, -0.87055909, -0.90188624, ..., -1.03067011,
-1.07636976, -0.50577476]])
```

```
In [62]: x = pd.DataFrame(x,columns = names)x

.head()
```

Out[62]:

RowNum	Customer	Geograp	Gende	Age	Tenu	Balanc	NumOfProd	HasCrCa	re e	IsActive	Me ber	Id	hy
r ucts rd mbe													
0	-1.731878	-0.783213	-	1.0959	0.29351		1.2258	-0.911583	0.646092			0.97024	
				1.0417			48						
			0.901886	88	7	60							
			-	-									
1	-1.731531	-0.606534	1.515067	1.0959	0.19816	1.3875		0.1173	-0.911583	-			
			0.97024	88	4	38	50	1.547768					
			-										
	-1.731185	-0.995885	-	1.0959	0.29351	1.0329	1.3330	2.527057	0.646092			-1.03067	
2			0.901886	88	7	08	53						
			-	-	-								
	-1.730838	0.144767	-	1.0959	0.00745	1.3875		1.2258	0.807737	-		-1.03067	
			0.901886	88	7	38	48	1.547768					
3			-	-								0.97024	
	-1.730492	0.652659	1.515067	1.0959	0.38887	1.0417	0.7857	0.911583	0.646092				
					88	1	60	28					

Split the data into training and testing

```
In [69]:

from sklearn.model_selection import train_test_split

# Split training and testing data

xtrain,xtest,ytrain,ytest = train_test_split(x,y,test_size=0.20,random_state=0)
```

In [70]:

```
# Checking shape of data
```

```
xtrain.shape,xtest.shape Out[70]: ((8000,  
12), (2000, 12))
```